


Dedicated at Thy Lotus Feet



# BlockChain Consensus & Smart HBasechainDB



Bangarugiri Sateesh • 18555

Supervisor : Dr. Pallav Kumar Baruah

# Contents


PROJECT INTERIM REVIEW

- Introduction
- Problem Statement
- Objectives
- Forking Resilient Kafka Orderer
- Smart HBasechainDB
- Future Work



# BLOCKCHAIN

# What is Blockchain ?




Distributed Ledger Technology  
(acronym: DLT)



Ledger records  
ALL transactions



'key/value' database with  
current state (optional)



Decentral peer-to-peer  
network of nodes




Public key-cryptography  
without central authority



Any transaction added is  
validated by multiple entities <sup>5</sup>

## Making a hash of it



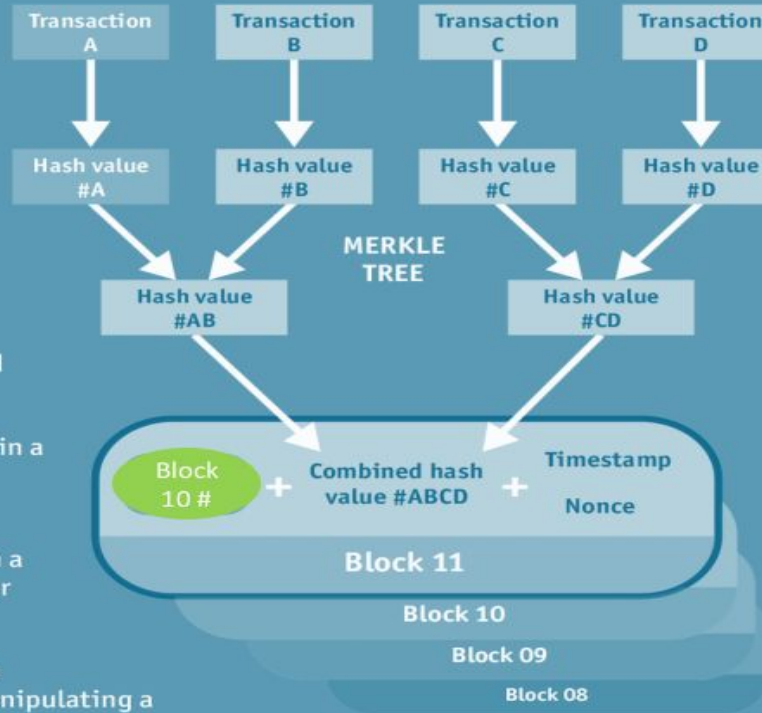
Each transaction in the set that makes up a block is fed through a program that creates an encrypted code known as the hash value.

Hash values are further combined in a system known as a Merkle Tree.

The result of all this hashing goes into the block's header, along with a hash of the previous block's header and a timestamp.

The header then becomes part of a cryptographic puzzle solved by manipulating a number called the nonce.

Once a solution is found the new block is added to the blockchain.



BUNDLING  
TRANSACTIONS  
INTO A BLOCK

CHAIN OF BLOCKS  
BY ADDING HASH OF  
THE PREVIOUS  
BLOCK TO CURRENT


TAMPER-PROOFING

# Proof of Work

The system is called **proof of work** because the probability of mining the block is increased with the amount of work that is put in.

**1** A very complex mathematical challenge is proposed to the blockchain network

**2** The miners have to compete to find the solution, which takes time and resources, making it costly for the contestants.



**3** The first miner to solve the problem has the ability to validate transactions and create a new block, receiving a reward afterwards.







# Ethereum



## Proof of Stake



In **Proof of Stake**, each validator owns some stake in the network, and has to lock it in order to be selected.

- 1** Anyone who holds the base cryptocurrency can become a validator, although sometimes a locked up deposit is required.




- 2** A validator's chance of mining a block is based on how much of a stake (or cryptocurrency) they have.

*For example, if you owned 1% of the cryptocurrency, you would be able to mine 1% of all its transactions.*




- 3** The PoS protocol will randomly assign the right to create a block in between selected validators, based upon the value of their stakes.

The chosen validator is rewarded by a part or the whole of the transaction fee.



# Types of Blockchains



# Public Blockchain Features



Decentralized



Transparency

## Public Blockchain

Allows anyone to participate



Rules immutable  
after being deployed



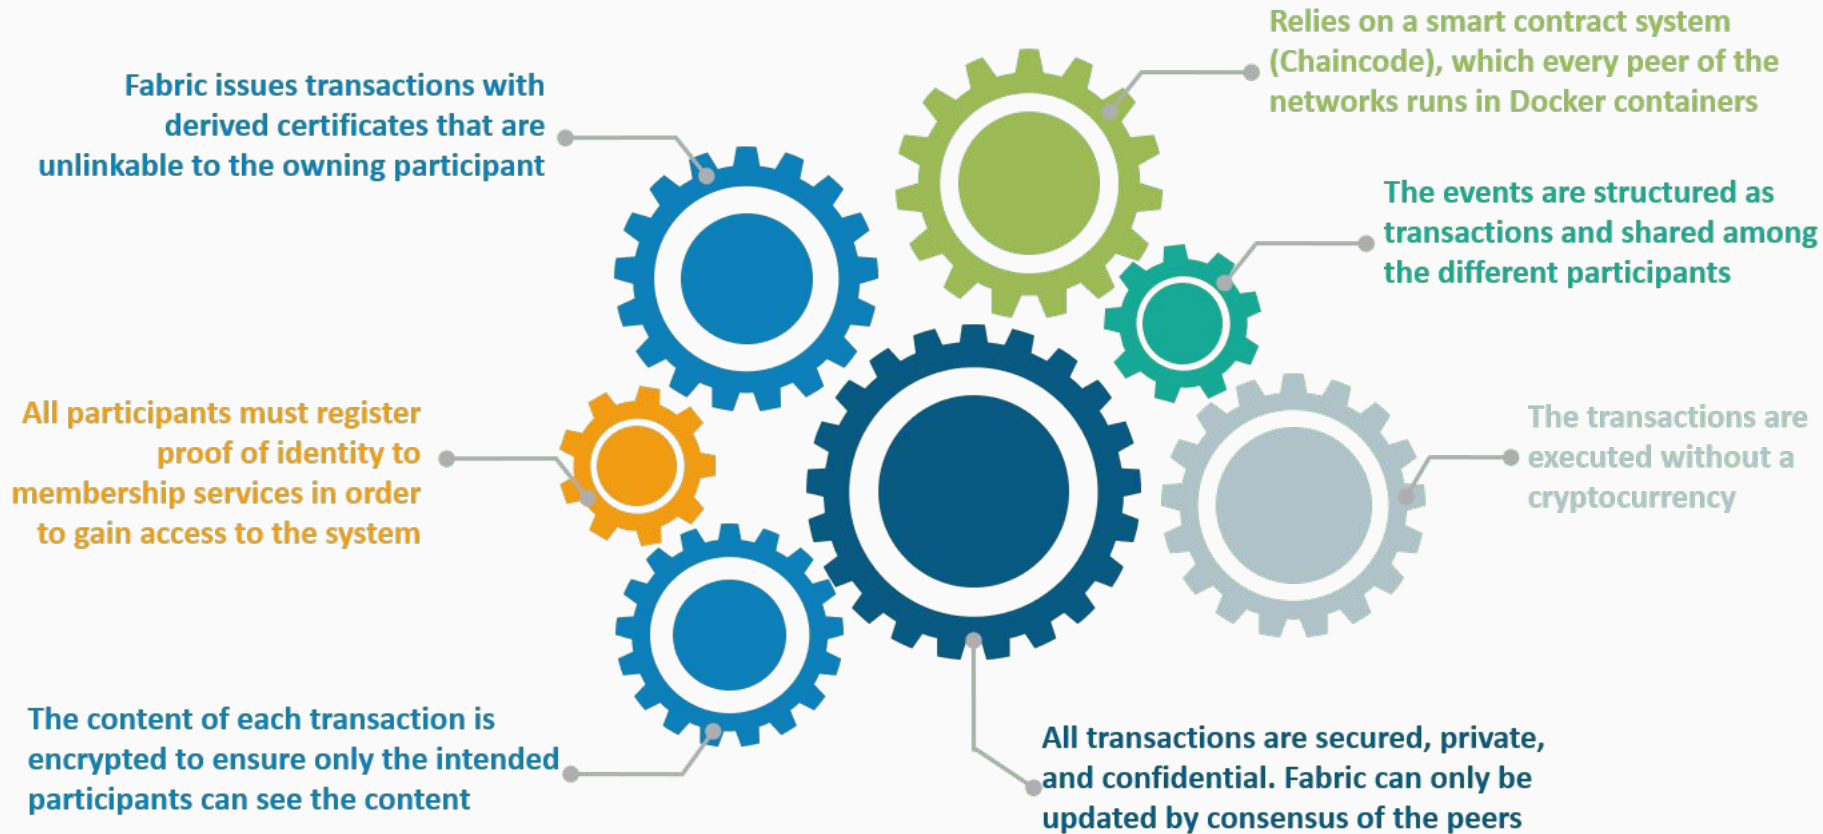
Validated by every  
participant



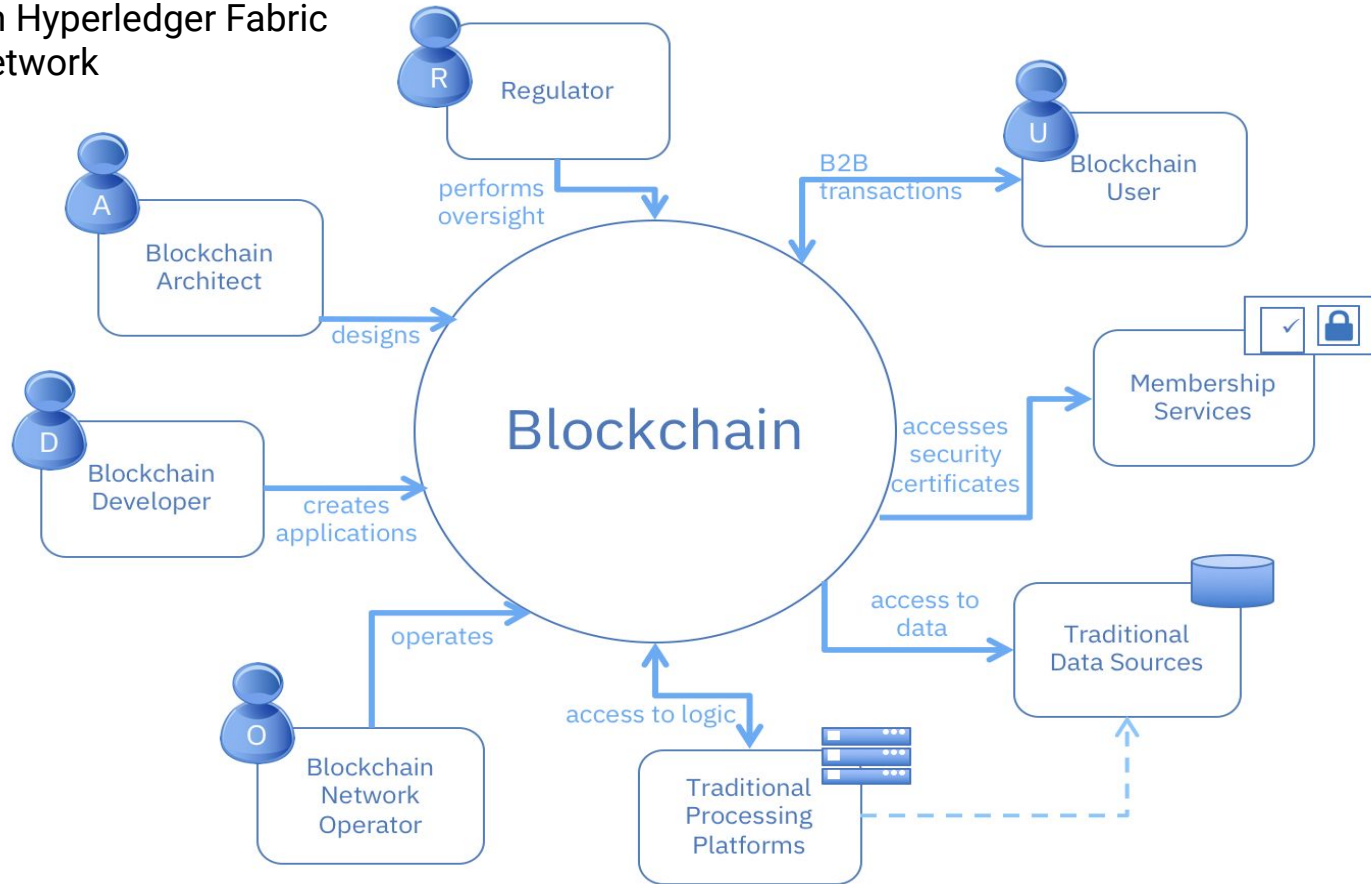
Slow



# **HYPERLEDGER** **FABRIC**

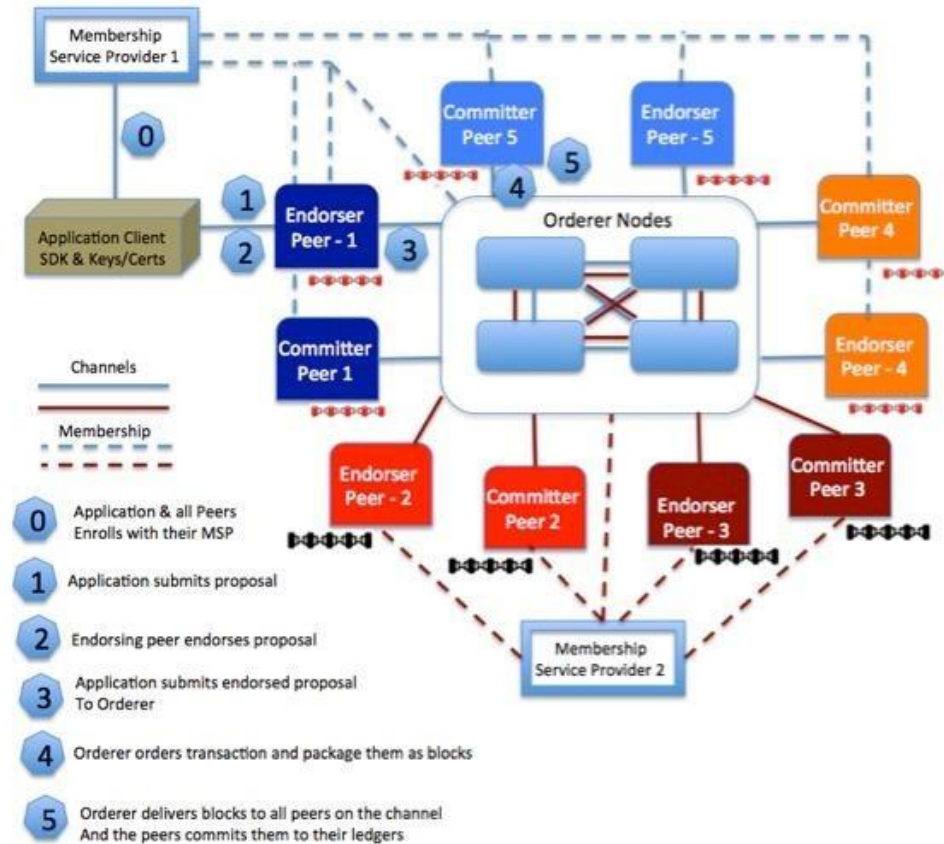


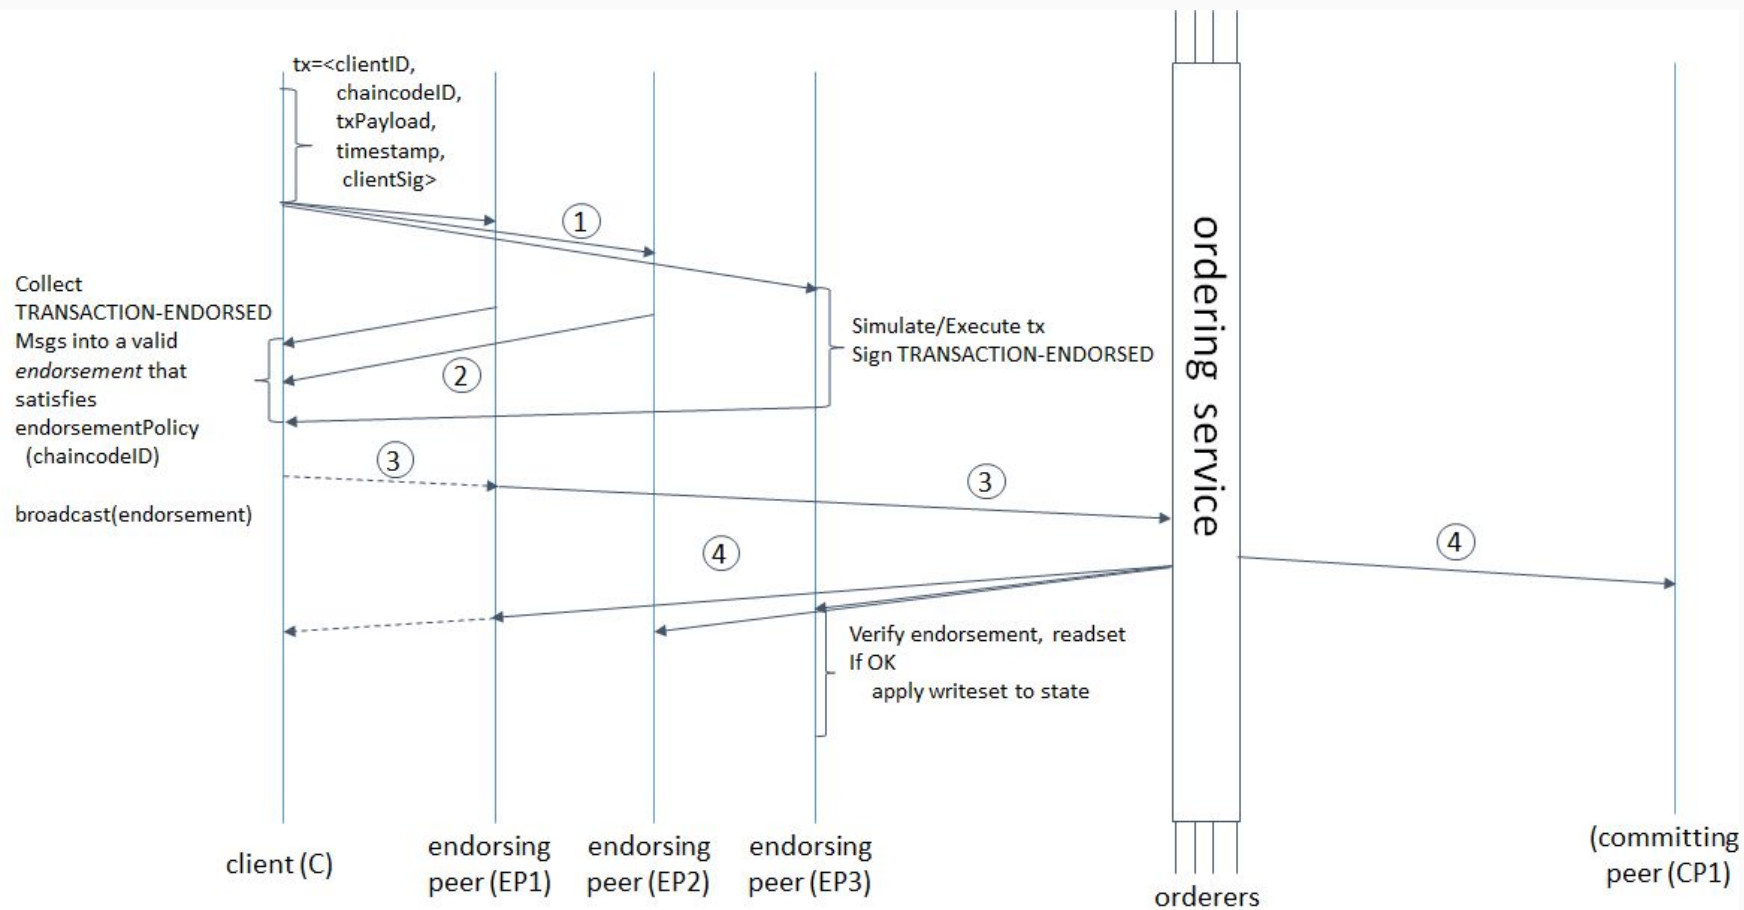
## Participants in Hyperledger Fabric Network



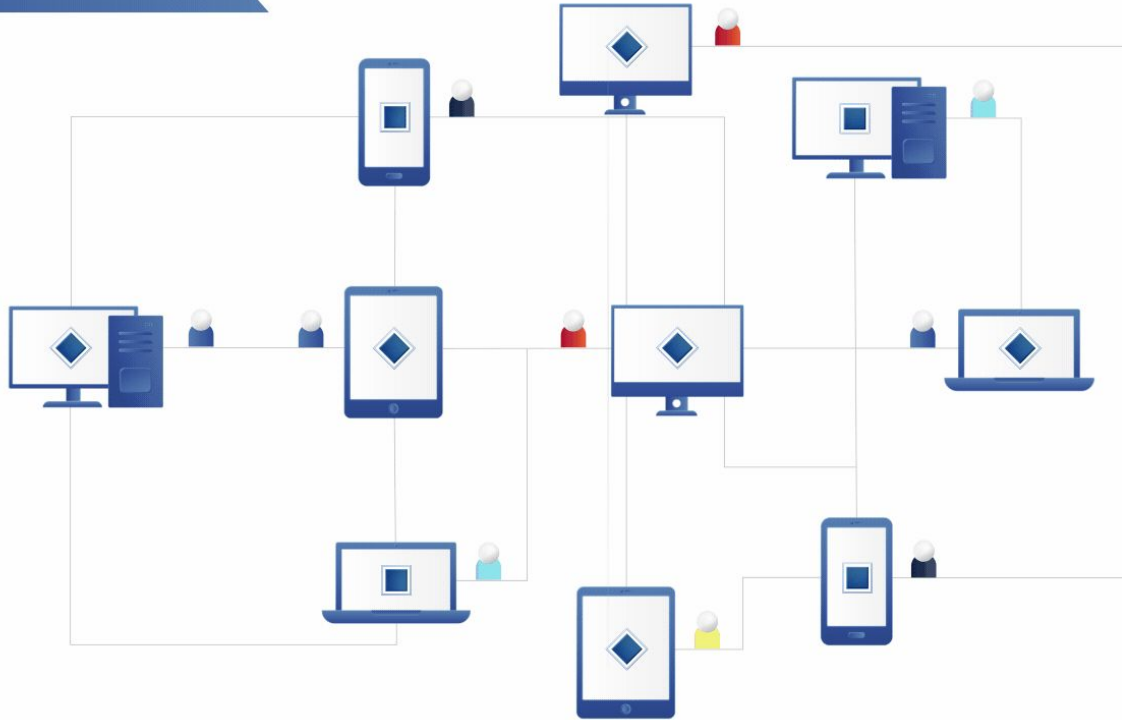


# HyperLedger Fabric 1.0 : Private Permissioned Blockchain





## Consensus Protocol



# Crash Fault Tolerance vs Byzantine Fault Tolerance

- **Crash Fault Tolerance (CFT)** is one level of resiliency, where the system can still correctly reach consensus if components fail. (eg. Kafka Orderer and Raft Orderer)
- Crash failure = the process halts
- Fail-Stop failure is a simple abstraction that mimics crash failure when process behavior becomes arbitrary.
- If a system cannot tolerate fail-stop failures, it cannot tolerate crash failures.
- **Byzantine Fault Tolerance (BFT)** is more complex and deals with systems that may have malicious actors.
- Numerous possible causes.
- Most difficult kind of failure to deal with.
- Some of the arbitrary node failures are given below : **Failure to return a result, Respond with an incorrect result, Respond with a deliberately misleading result, Respond with a different result to different parts of the system.**

# Problem Statement

To build a forking resilient Kafka ordering service by modifying existing one and ensuring safety guarantees.

To make existing HBasechainDB support Smart Contract by modifying its architecture and check for flexibility.

# Objectives

- ❑ To study the Bitcoin blockchain technology for zeroing on problems to work upon
- ❑ To explore newer blockchain implementations – Ethereum and Hyperledger
- ❑ To write smart contracts in the Ethereum framework and acquire an understanding of its use cases
- ❑ To study approaches taken to addressing Byzantine Faults in blockchain
- ❑ To explore consensus algorithms like Paxos, Raft, Kafka Orderer, Honey Badger
- ❑ To study the HBasechainDB1.0 and HBasechainDB2.0 for understanding its architecture
- ❑ To implement modified Kafka Orderer for a forking resilience in Hyperledger Fabric in GO language
- ❑ To accommodate the changes in architecture of HBasechainDB2.0 for supporting Smart Contract and check for flexibility



# Work Done

- ✓ To study the Bitcoin blockchain technology for zeroing on problems to work upon
- ✓ To explore newer blockchain implementations – Ethereum and Hyperledger
- ✓ To write smart contracts in the Ethereum framework and acquire an understanding of its use cases
- ✓ To study approaches taken to addressing Byzantine Faults in blockchain
- ✓ To explore consensus algorithms like Paxos, Raft, Kafka Orderer, Honey Badger
- ✓ To study the HBasechainDB1.0 and HBasechainDB2.0 for understanding its architecture
- ❑ To implement modified Kafka Orderer for a forking resilience in Hyperledger Fabric in GO language
- ❑ To accommodate the changes in architecture of HBasechainDB2.0 for supporting Smart Contract and check for flexibility

# Literature

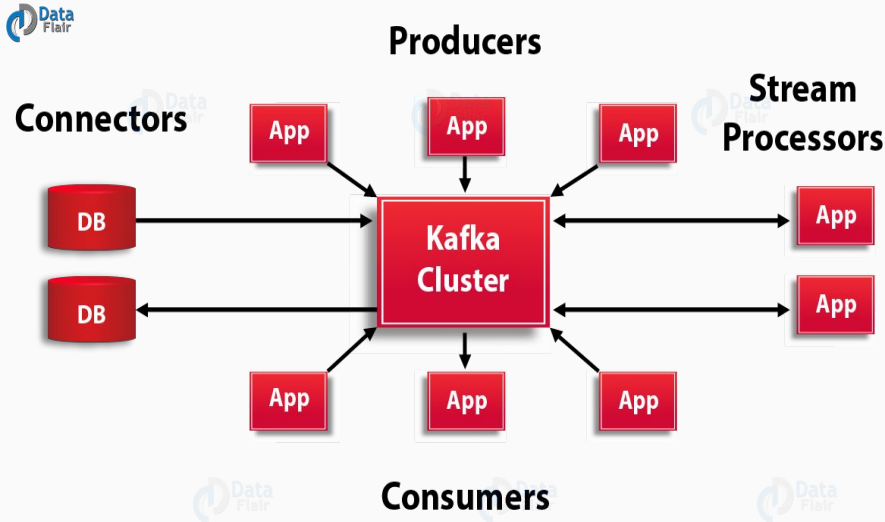
- Elli Androulaki, Artem Barger, Vita Bortnikov, Christian Cachin, Konstantinos Christidis, Angelo De Caro, David Enyeart, Christopher Ferris, Gennady Laventman, Yacov Manevich, Manish Sethi, Gari Singh, Keith Smith, Alessandro Sorniotti, Chrysoula Stathakopoulou, Marko Vukolić, Sharon Weed, Cocco, Jason Yellick. (2018). "Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains".
- Christian Gorenflo, Stephen Lee, Lukasz Golab, S. Keshav, (March 2019). "FastFabric: Scaling Hyperledger Fabric to 20,000 Transactions per Second".
- Arati Baliga, Nitesh Solanki, Shubham Verekar, Amol Pednekar, Pandurang Kamat and Siddhartha Chatterjee, "Performance Characterization of Hyperledger Fabric", 2018 Crypto Valley Conference on Blockchain Technology.
- Diego Ongaro and John Ousterhout, (May 2013). "In Search of an Understandable Consensus Algorithm"
- Leslie Lamport, (01 Nov 2001). "Paxos Made Simple"
- Medium Blogs



Distributed publish subscribe  
messaging system

- It was developed at LinkedIn and later on became a part of Apache Project.
- It is fast, agile, scalable and distributed by design.
- A message handling system, that uses the Publish-Subscribe Model.
- Consumers subscribe to Topic to receive new messages , that are published by Producer.

# Kafka Architecture and Terminology




- Topic : A stream of messages belonging to a particular category is called a topic
- Producer : A producer can be any application that can publish messages to a topic
- Consumer : A consumer can be any application that subscribes to topics and consumes the messages
- Broker : Kafka cluster is a set of servers, each of which is called a broker

# Kafka Orderer in Hyperledger Fabric

- For every channel, we have a separate partition.
- Each channel maps to the single partition topic.
- The OSNs after confirming the permissions in the chain relay the incoming client transactions (received via the Broadcast RPC)
- The OSNs can then consume that partition and get back an ordered list of transactions that is common across all ordering service nodes.
- The transactions in a chain are batched, with a timer service. That is, whenever the first transaction for a new batch comes in, a timer is set.
- The block (batch) is cut either when the maximum number of transactions are reached (or when the timer expires, whichever comes first).
- Each OSN maintains a local log for every chain, and the resulting blocks are stored in a local ledger.
- The transaction blocks are then served to the receiving clients via the Deliver RPC.
- In the case of a crash, the relays can be sent through a different OSN since all the OSNs maintain a local log. This has to be explicitly defined, however.

# Forking Resilient Kafka Orderer


- Blockchain Fork: Occurs when two or more blocks have the **same block height** and the block height of a particular block is defined as the number of blocks preceding it in the blockchain.
- Fork happens when two miners find a block at nearly the same time in Bitcoin blockchain.
- Forking in each and every blockchain is different, based on the design architecture.
- Fork in Hyperledger Fabric:



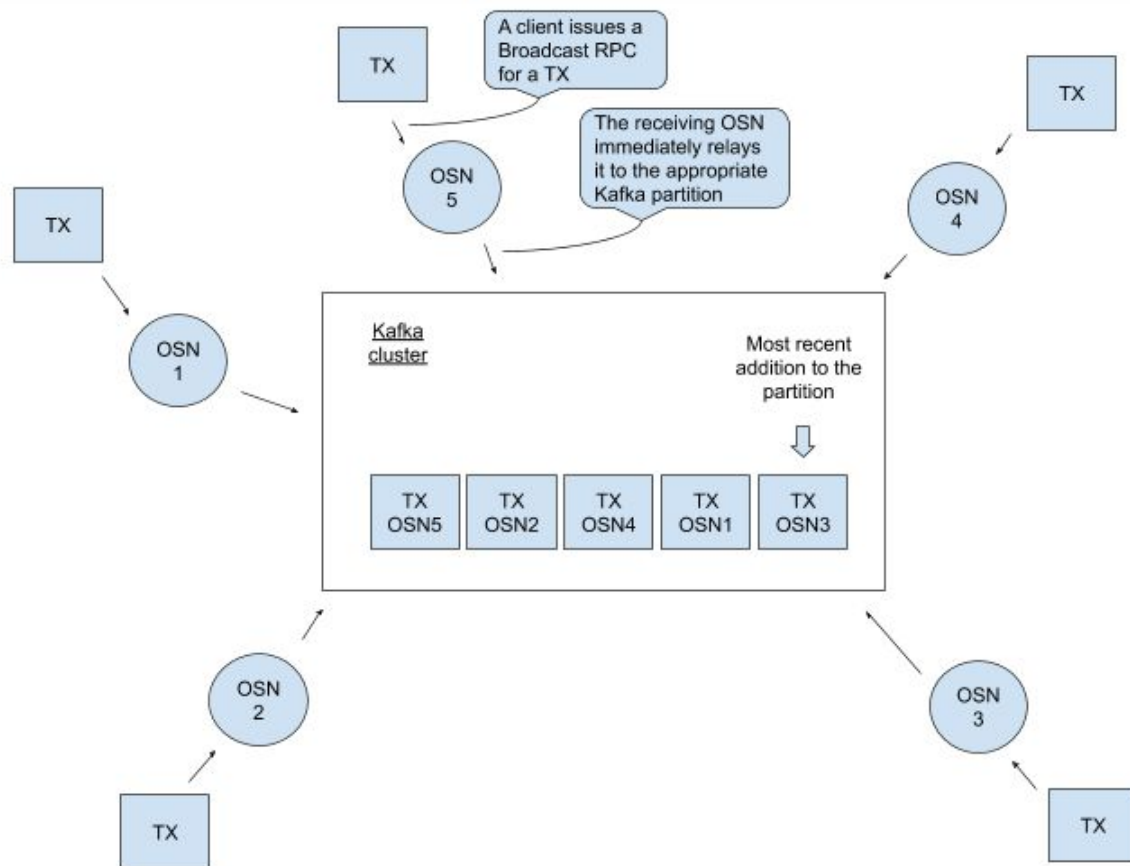
Source[VI]



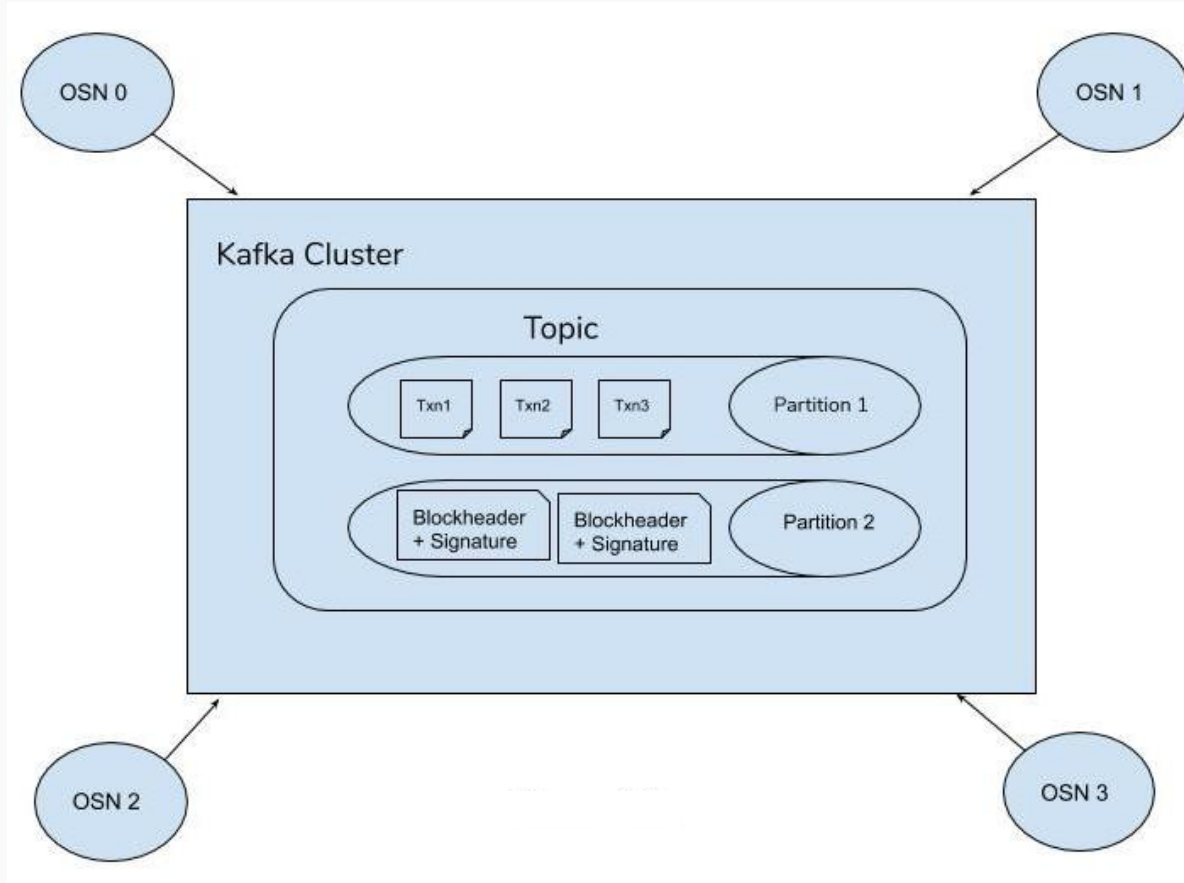
# Kafka Ordering Service in Hyperledger Fabric



# Kafka Ordering Service in Hyperledger Fabric



# Modified Kafka Ordering Service in Hyperledger Fabric for Forking Resilience



# Steps to avoid Forking in Kafka Orderer

- Create a second partition in the kafka topic for each channel apart from the first partition from where the OSNs consume the transactions.
- Have each OSN produce its signature over the block header, but before committing it, produce the signature and block header into the second partition.
- Each OSN commit their block into their local logs after getting a sufficient number of matching signatures from the second partition and all these signatures must be added into block metadata.



# HBase

open source, multidimensional, distributed,  
scalable and a *NoSQL database*

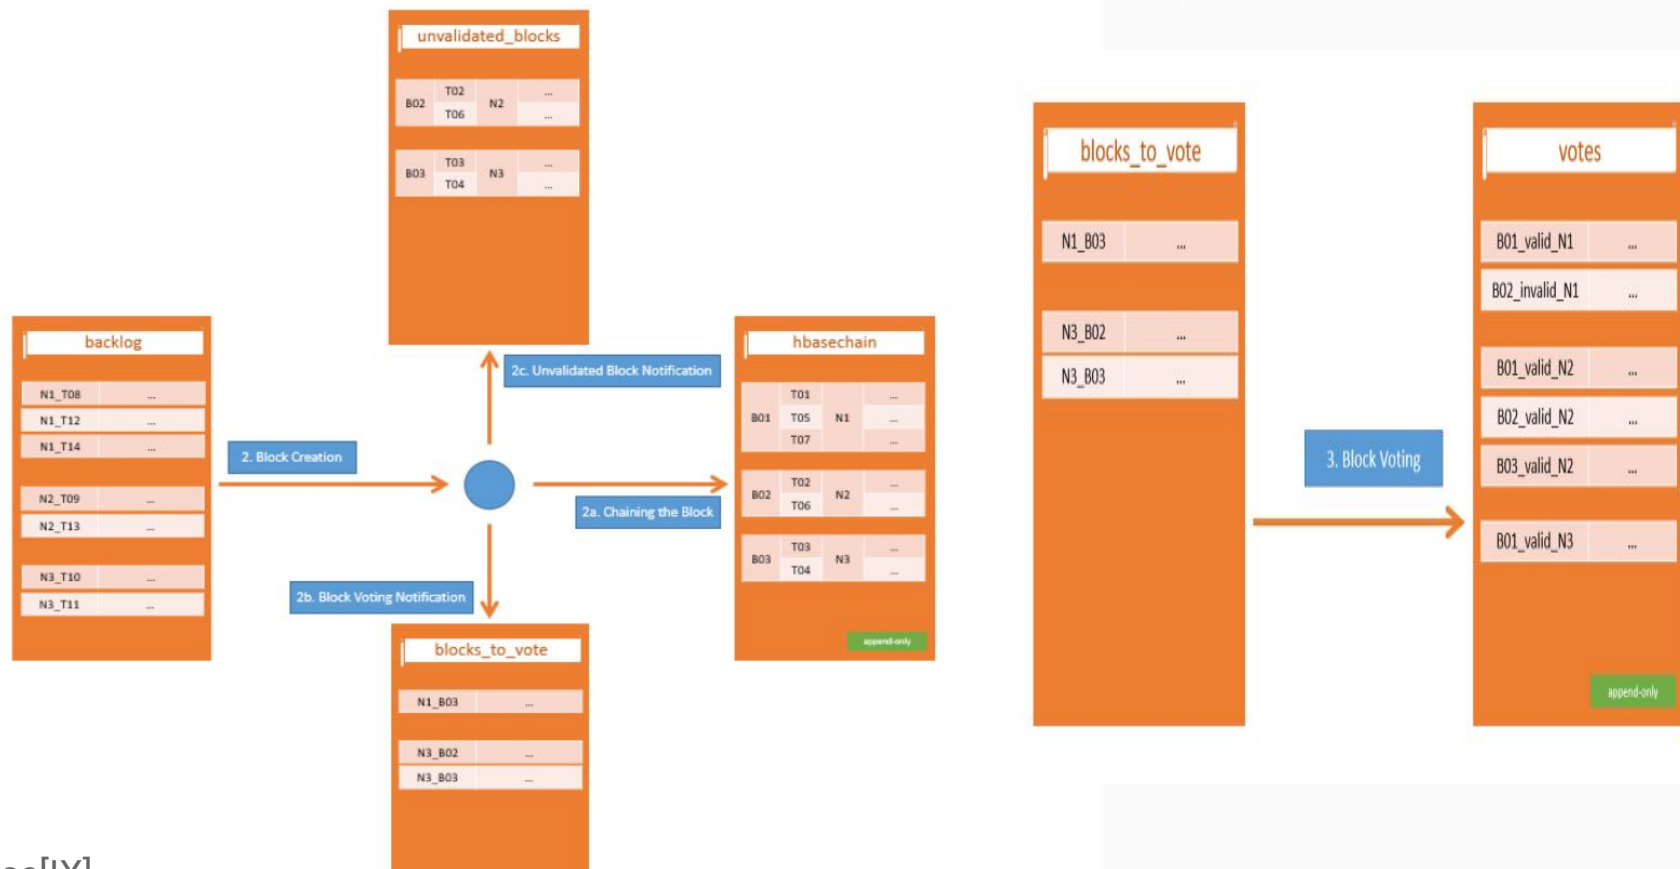
- It runs on top of HDFS(Hadoop Distributed File System)
- Designed to provide a fault tolerant way of storing large collection of sparse datasets.
- It achieves high throughput and low latency by providing faster Read/Write Access on huge datasets.
- It is the choice for the applications which require fast & random access to large amounts of data.

# HBasechainDB

- A scalable blockchain framework on Hadoop ecosystem.
- It supports basic facilities like CREATION and TRANSFER of asset.
- It is operated using Federation of Nodes. All the nodes in the federation have equal privileges which gives HBasechainDB its decentralization.
- Any client can submit or retrieve transactions or blocks, but only the federation nodes can modify the blockchain.
- HBasechainDB's transaction model consists of; Transaction Id, Asset, List of Inputs, List of Outputs and Metadata.
- The current implementation of HBasechainDB uses six HBase tables: backLog, block, hbasechaindb, toVote, vote, reference.
- Operation of HBasechainDB involves a sequence of the following three steps: Transaction Insertion, Block Creation, and Block Voting.



# Block Creation and Voting



# HBasechainDB

HBasechainDB1.0  
**Adarsh Saraf**

HBasechainDB2.0  
**Subhankar Sahoo**

# Smart HBasechainDB

HBasechainDB  
**Adars**

HBasechainDB3.0  
**Sateesh**

HBasechainDB2.0  
**r Sahoo**




# Timeline

Studied Blockchain basics(Bitcoin), Ethereum, Certification courses on Blockchain Technology.

Writing Chaincodes(in Fabric), Finalized Problem Statement, Studied Consensus algorithms(in Blockchain), Identifying Risk areas in Kafka Orderer

Installing HBase on 8-node Hadoop Cluster and HBasechainDB



Writing Smart Contracts in Ethereum, Learnt Hyperledger Fabric System, Network setup for Fabric, Explored Cosmos Network and Edge Computing Platforms

Identified Fork situation in Kafka Order, Modifying Kafka Orderer, Studied HBasechainDB Architecture, Writing HiPC Paper on Forking Resilience in Kafka Orderer.

# Future Work

## Step 1

Installing HBase on 8-node Hadoop Cluster and HBasechainDB

## Step 2

To implement modified Kafka Orderer for a forking resilience in Hyperledger Fabric in GO language and check for performance

## Step 3

To accommodate the changes in architecture of HBasechainDB2.0 for supporting Smart Contract and check for flexibility

# Acknowledgements

- Dr. Pallav Kumar Baruah
- Sri. Krishnamoorthy Balaraman
- Sri. Varun Vallampatla
- Jason Yellick, Blockchain Software Developer at IBM, North-Carolina Area, USA
- Soumya Nayak, System Engineer, First American India Pvt Ltd, Hyderabad
- Subhankar Sahoo

# Image References

1. <https://www.1point21gws.com/insights/blockchain/enterprise-cloud-security-is-blockchain-technology-the-missing-link/>
2. “Build a decentralized blockchain application” by Robert van Mölken Oracle Developer ChampionBlockchain / Integration Specialist (AMIS)
3. “Blockchain-based Bigdata Solutions” Interim presentation by Adarsh Saraf
4. <https://www.cryptolinenews.com/top-stories/facebook-plans-to-launch-its-own-cryptocurrency/>
5. <https://lisk.io/academy/blockchain-basics/how-does-blockchain-work/proof-of-work>
6. <https://acadgild.com/blog/types-of-blockchain>
7. [https://medium.com/@yoa\\_project/types-of-blockchains-fc94ef840043](https://medium.com/@yoa_project/types-of-blockchains-fc94ef840043)
8. <https://medium.com/coinmonks/beginning-with-the-hyperledger-fabric-241b54859476>
9. <https://www.edureka.co/blog/hyperledger-fabric/>
10. <https://medium.com/predict/stakeholders-in-a-blockchain-solution-d54591844c20>
11. “Blockchain-based Bigdata Solutions” Interim presentation by Adarsh Saraf



# Image References (Cont.)

- I. <http://websecuritypatterns.com/blogs/2017/07/07/unpacking-hyperledger-fabric-1-0-under-the-hood-of-a-permissioned-blockchain/>
- II. <https://hyperledger-fabric.readthedocs.io/en/release-1.4/arch-deep-dive.html#swimlane>
- III. <https://www.quora.com/What-is-consensus-in-blockchain>
- IV. <https://www.shutterstock.com/search/coche+valid%C3%A9>
- V. <https://kafka.apache.org/>
- VI. “Adding forking resilience in Kafka orderer even in the presence of Byzantine Faults” by Bangarugiri Sateesh, Pallav Kumar Baruah, Jason Yellick
- VII. <https://docs.google.com/document/d/19JihmW-8bITzN99IAubOfseLUZqdrB6sBR0HsRgCAnY/edit>
- VIII. <https://docs.google.com/document/d/19JihmW-8bITzN99IAubOfseLUZqdrB6sBR0HsRgCAnY/edit>
- IX. “HBaseChainDB2.0” by Subhankar Sahoo and Pallav Kumar Baruah
- X. <https://gramboard.ai/blog/get-instagram-explore-page-2019/>

# References

- <https://www.edureka.co/blog/hbase-tutorial>
- <https://www.geeksforgeeks.org/practical-byzantine-fault-tolerancepbft/>
- <https://codeburst.io/the-abcs-of-kafka-in-hyperledger-fabric-81e6dc18da56>
- <https://fabrictestdocs.readthedocs.io/en/latest/glossary.html>
- <http://homepage.divms.uiowa.edu/~ghosh/16612.week10.pdf>
- <https://www.cs.princeton.edu/courses/archive/fall16/cos418/docs/L9-bft.pdf>
- <https://www.investopedia.com/terms/b/block-height.asp>
- DataFlair, Edureka.

